



Experiment No: 05

Code 01 :-

```
import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

# Import warnings

import warnings

warnings.filterwarnings("ignore")

# We will use some methods from the sklearn module

from sklearn import linear_model

from sklearn.linear_model import LinearRegression

from sklearn import metrics

from sklearn.metrics import mean_squared_error, mean_absolute_error

from sklearn.model_selection import train_test_split, cross_val_score

# Load the dataset

df = pd.read_csv("Student_Performance.csv")

print(df.corr(numeric_only=True)) # Use numeric_only=True to only get numeric
correlations

print(df.describe())

# Setting the value for X and y

X = df[['Hours_Studied', 'Previous_Scores']] # Adjusted features
```



```
y = df['Performance_Index'] # Adjusted target variable

# Visualizing the features

fig, axs = plt.subplots(2, figsize=(5, 10)) # Adjusted height for better spacing

plt1 = sns.boxplot(x=df['Hours_Studied'], ax=axs[0]) # Adjusted for the new feature

plt2 = sns.boxplot(x=df['Previous_Scores'], ax=axs[1]) # Adjusted for the new feature

plt.tight_layout()

sns.histplot(df['Performance_Index'], kde=True) # Changed to histplot with kde for better
visualization

sns.pairplot(df, x_vars=['Hours_Studied', 'Previous_Scores'], y_vars='Performance_Index',
height=4, aspect=1, kind='scatter')

plt.show()

# Create the correlation matrix and represent it as a heatmap.

sns.heatmap(df.corr(numeric_only=True), annot=True, cmap='coolwarm') # Use
numeric_only=True for heatmap

plt.show()

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=100)

# Fitting the Multiple Linear Regression model

reg_model = LinearRegression().fit(X_train, y_train)

# Printing the model coefficients

print('Intercept: ', reg_model.intercept_)
```

```
# Pair the feature names with the coefficients

print(list(zip(X.columns, reg_model.coef_))) # Corrected to use .columns for feature names


# Predicting the Test and Train set result

y_pred = reg_model.predict(X_test)

x_pred = reg_model.predict(X_train) # This is actually predicting on the train set


print("Prediction for test set: {}".format(y_pred))


# Actual value and the predicted value

reg_model_diff = pd.DataFrame({'Actual value': y_test, 'Predicted value': y_pred})

print(reg_model_diff) # Added print to display the DataFrame


mae = metrics.mean_absolute_error(y_test, y_pred)

mse = metrics.mean_squared_error(y_test, y_pred)

r2 = np.sqrt(mse) # Changed to use mse instead of mean_squared_error directly


print('Mean Absolute Error:', mae)

print('Mean Square Error:', mse)

print('Root Mean Square Error:', r2)
```

Code 01 Output :-



NUTAN MAHARASHTRA VIDYA PRASARAK MANDAL'S
NUTAN COLLEGE OF ENGINEERING & RESEARCH (NCER)

DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING - ARTIFICIAL INTELLIGENCE



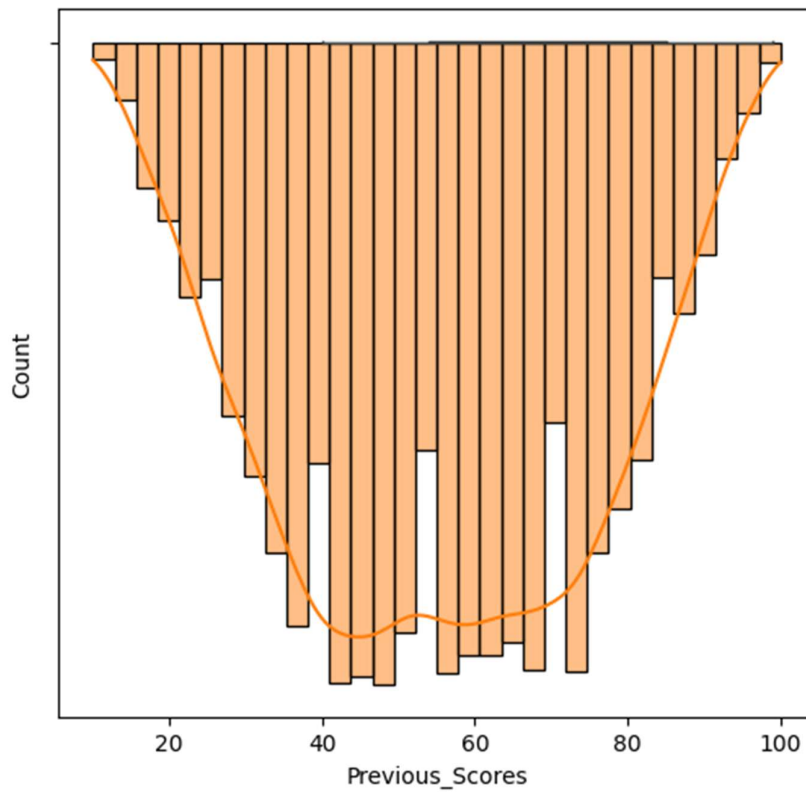

```
Hours_Studied      1.000000      -0.012390      0.001245 \
Previous_Scores    -0.012390      1.000000      0.005944
Sleep_Hours        0.001245      0.005944      1.000000
Sample_Question_Papers_Practiced  0.017463      0.007888      0.003990
Performance_Index   0.373730      0.915189      0.048106
```

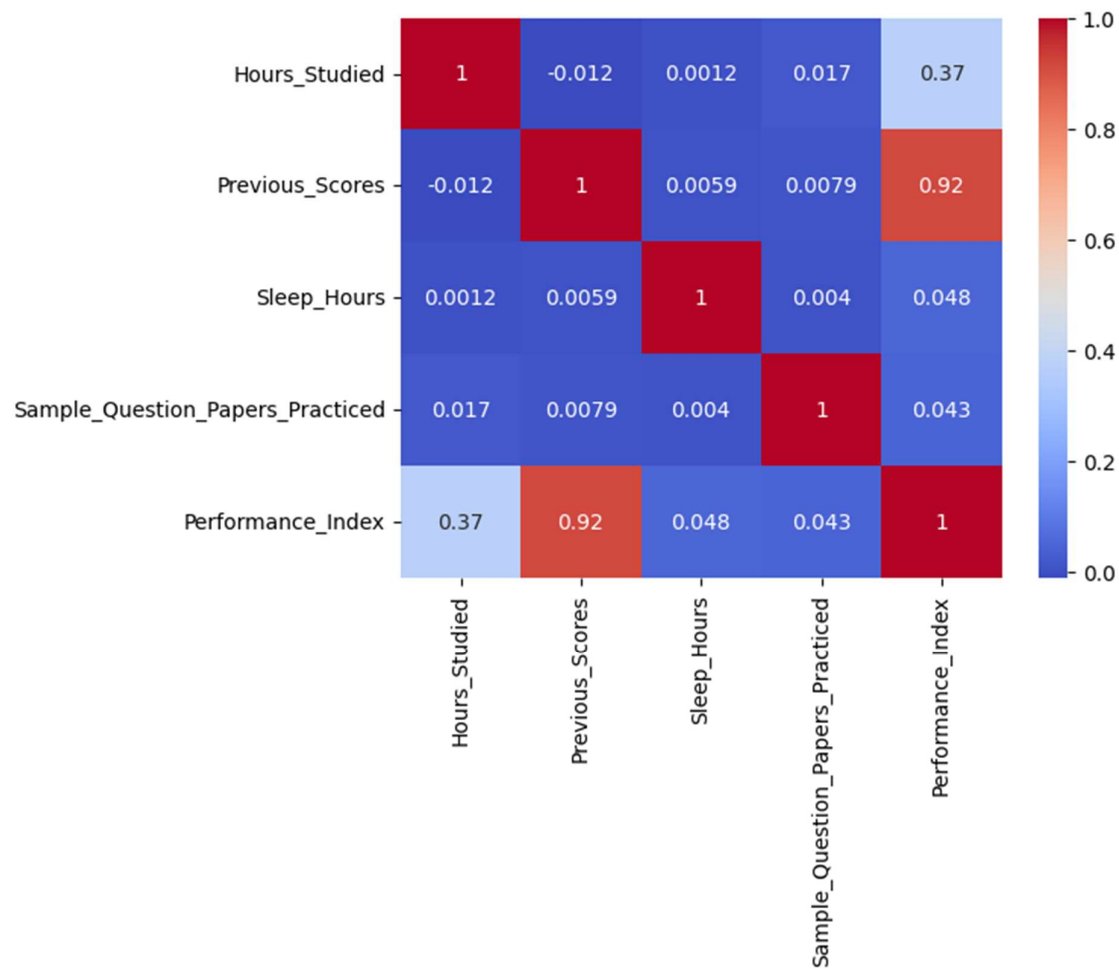
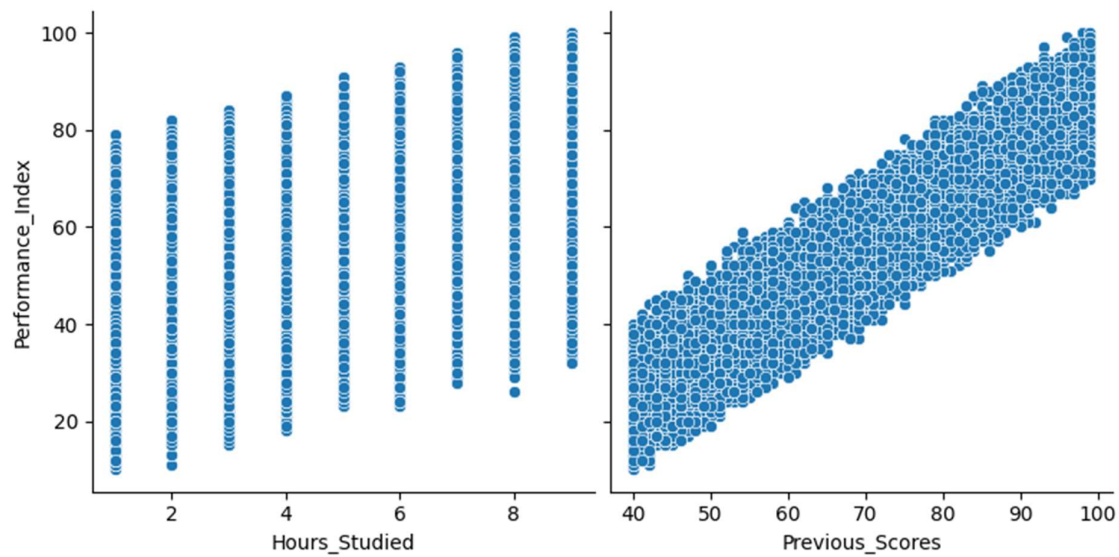
```
Sample_Question_Papers_Practiced \
Hours_Studied      0.017463
Previous_Scores     0.007888
Sleep_Hours        0.003990
Sample_Question_Papers_Practiced  1.000000
Performance_Index   0.043268
```

```
Performance_Index
Hours_Studied      0.373730
Previous_Scores     0.915189
Sleep_Hours        0.048106
Sample_Question_Papers_Practiced  0.043268
Performance_Index   1.000000
```

```
Hours_Studied      Previous_Scores      Sleep_Hours \
count  10000.000000  10000.000000  10000.000000
mean    4.992900      69.445700      6.530600
std     2.589309      17.343152      1.695863
min     1.000000      40.000000      4.000000
25%     3.000000      54.000000      5.000000
50%     5.000000      69.000000      7.000000
75%     7.000000      85.000000      8.000000
max     9.000000      99.000000      9.000000
```

```
Sample_Question_Papers_Practiced      Performance_Index
count  10000.000000  10000.000000
mean    4.583300      55.224800
std     2.867348      19.212558
min     0.000000      10.000000
25%     2.000000      40.000000
50%     5.000000      55.000000
75%     7.000000      71.000000
max     9.000000      100.000000
```





```
Intercept: -29.759696031285607  
[('Hours_Studied', np.float64(2.8458785973597283)), ('Previous_Scores', np.float64(1.0192750395624741))]  
Prediction for test set: [65.62826464 49.70329949 34.41417389 ... 71.9558614 35.22150241  
14.87645919]
```

	Actual value	Predicted value
8018	65	65.628265
9225	51	49.703299
3854	35	34.414174
2029	97	92.896287
3539	87	84.530140
...
4208	69	67.242922
5335	60	58.705286
2554	73	71.955861
13	33	35.221502
4468	15	14.876459

```
[3000 rows x 2 columns]  
Mean Absolute Error: 1.8372246647574142  
Mean Square Error: 5.3409802420571  
Root Mean Square Error: 2.3110560880379127
```

Code 02 :-

```
# Import necessary libraries  
  
import pandas as pd  
  
import numpy as np  
  
import matplotlib.pyplot as plt  
  
import seaborn as sns  
  
from sklearn import linear_model  
  
from sklearn.model_selection import train_test_split  
  
from sklearn import metrics  
  
  
# Load the dataset  
  
df = pd.read_csv("Student_Performance.csv")  
  
  
  
# Prepare the features and target variable  
  
X = df[['Hours_Studied', 'Previous_Scores']] # Adjusted features
```



```
y = df['Performance_Index'] # Adjusted target variable

# Perform exploratory data analysis

print(df.corr(numeric_only=True)) # Only get numeric correlations

sns.pairplot(df, x_vars=['Hours_Studied', 'Previous_Scores'], y_vars='Performance_Index',
height=4, aspect=1, kind='scatter')

plt.show()

# Split the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Create and fit the multivariable regression model

reg_model = linear_model.LinearRegression()

reg_model.fit(X_train, y_train)

# Print model coefficients

print('Intercept:', reg_model.intercept_)

print('Coefficients:', reg_model.coef_)

# Make predictions

y_pred = reg_model.predict(X_test)

# Evaluate the model

mae = metrics.mean_absolute_error(y_test, y_pred)

mse = metrics.mean_squared_error(y_test, y_pred)

r2 = metrics.r2_score(y_test, y_pred)
```



```
print('Mean Absolute Error:', mae)

print('Mean Squared Error:', mse)

print('R-squared:', r2)


# Visualize actual vs predicted values

plt.scatter(y_test, y_pred)

plt.xlabel('Actual Performance Index')

plt.ylabel('Predicted Performance Index')

plt.title('Actual vs Predicted Performance Index')

plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--') # 45-degree line


plt.show()
```

Code 02 Output: -

```
Hours_Studied      1.000000      -0.012390      0.001245 \
Previous_Scores    -0.012390      1.000000      0.005944
Sleep_Hours        0.001245      0.005944      1.000000
Sample_Question_Papers_Practiced  0.017463      0.007888      0.003990
Performance_Index  0.373730      0.915189      0.048106


Sample_Question_Papers_Practiced \
Hours_Studied      0.017463
Previous_Scores    0.007888
Sleep_Hours        0.003990
Sample_Question_Papers_Practiced  1.000000
Performance_Index  0.043268


Performance Index
Hours_Studied      0.373730
Previous_Scores    0.915189
Sleep_Hours        0.048106
Sample_Question_Papers_Practiced  0.043268
Performance_Index  1.000000
```



Intercept: -29.60536090045995
Coefficients: [2.86449217 1.01615225]
Mean Absolute Error: 1.8101437492369952
Mean Squared Error: 5.201346140631689
R-squared: 0.9859445764337783

